



AMERICAN INTERNATIONAL UNIVERSITY-BANGLADESH

MID-TERM PROJECT

Course: INTRODUCTION TO DATA SCIENCE

Sec: A

SUBMITTED TO

Faculty: Tohedul Islam

SUBMITTED BY

Group: 15

NAME	ID
1. Abdullah Muhammad Hamja	20 -43465-1
2. Sumaiya Ahmed Susmita	21-45266-2

Heart Disease Classification Dataset

Cardiovascular diseases (CVDs) stand as a leading global cause of mortality, encompassing conditions like coronary heart disease, cerebrovascular disease, and more. With 17.9 million annual deaths attributed to CVDs, understanding and predicting factors contributing to heart attacks are crucial for public health. This dataset comprises 1319 samples, each characterized by eight input fields—age, gender, heart rate, systolic and diastolic blood pressure, blood sugar level, CK-MB, and Test-Troponin. The ninth field, 'class,' serves as the output, indicating the presence (positive) or absence (negative) of a heart attack. The dataset serves as a valuable resource for developing models to predict and understand the risk factors associated with cardiovascular events, contributing to advancements in preventive healthcare strategies.

Import CSV file in RStudio:

Code:

```
dataSet<-read.csv("D:/data.csv",header=TRUE,sep=",")
```

Dataset

The screenshot shows the RStudio interface. The top pane displays the source code for a script named 'dataSet'. The code consists of two lines: `dataSet<-read.csv("D:/data.csv",header=TRUE,sep=",")` and `dataSet`. The bottom pane shows the console output, which includes the R version (4.3.1) and the execution of the code. The output displays the structure of the dataset and its contents as a table with 28 rows and 8 columns: age, gender, impluse, pressurehigh, pressurelow, glucose, and class. The 'class' column contains values 'negative' and 'positive'.

```
R 4.3.1 · ~/
> dataSet<-read.csv("D:/data.csv",header=TRUE,sep=",")
> dataSet
```

	age	gender	impluse	pressurehigh	pressurelow	glucose	class
1	64	male	66	160	83	160	negative
2	21	male	94	98	46	296	positive
3	55	male	64	-160	77	270	negative
4	64	male	70	120	55	270	positive
5	55	male	64	112	65	300	negative
6	58	female	61	112	58	87	negative
7	32	female	40	179	68	102	negative
8	63	male	60	214	82	87	positive
9	44	female	60	NA	81	135	negative
10	67		61	160	95	100	negative
11	NA	female	60	166	90	102	negative
12	63	female	60	150	10	198	negative
13	64	male	60	199	5	92	positive
14	54	female	94	122	67	97	negative
15	47	male	76	120	70	319	negative
16	61	male	81	NA	66	134	positive
17	86	female	73	114	68	87	positive
18	45	female	70	100	68	96	negative
19	37	female	72	107	86	274	negative
20	45	male	60	109	65	89	positive
21	60	male	92	151	78	301	negative
22	48	male	135	98	60	100	positive
23	52	male	76	109	85	227	positive
24	30	male	63	110	68	107	positive
25	NA	male	63	320	63	269	positive
26	72	male	64	106	68	111	positive
27	42	male	65	150	68	101	negative
28	72	female	64	325	60	95	negative

Summary of the dataset: It's used to show a summary of the dataset.

Code: `summary(dataSet)`

```
Console Terminal x Background Jobs x
R 4.3.1 ~ /
> summary(dataSet)
      age      gender      impluse      pressureheight
Min.   : 19.00   Length:150   Min.    : 40.00   Min.    :-160.0
1st Qu.: 46.00   Class :character 1st Qu.: 62.00   1st Qu.: 110.0
Median : 56.00   Mode  :character Median : 74.00   Median : 121.5
Mean   : 56.14                      Mean   : 81.98   Mean   : 127.1
3rd Qu.: 64.00                      3rd Qu.: 83.00   3rd Qu.: 138.5
Max.   :155.00                      Max.    :1111.00 Max.    : 325.0
NA's   : 5                                NA's    : 2

      pressurelow      glucose      class
Min.   : 5.00   Min.    : 66.00   Length:150
1st Qu.:60.25   1st Qu.: 97.25   Class :character
Median :69.00   Median :116.00   Mode  :character
Mean   :68.95   Mean   :148.65
3rd Qu.:80.00   3rd Qu.:179.25
Max.   :95.00   Max.    :392.00

> |
```

Structure of the dataset: we can use `str` function to display the whole structure of the dataset.

Code: `str(dataSet)`

```
Console Terminal x Background Jobs x
R 4.3.1 ~ /
> str(dataSet)
'data.frame': 150 obs. of 7 variables:
 $ age      : int  64 21 55 64 55 58 32 63 44 67 ...
 $ gender    : chr  "male" "male" "male" "male" ...
 $ impluse   : int  66 94 64 70 64 61 40 60 60 61 ...
 $ pressureheight: int 160 98 -160 120 112 112 179 214 NA 160 ...
 $ pressurelow  : int  83 46 77 55 65 58 68 82 81 95 ...
 $ glucose     : int 160 296 270 270 300 87 102 87 135 100 ...
 $ class       : chr  "negative" "positive" "negative" "positive" ...

>
> |
```

Displaying Dataset: Ctrl + left-click (mouse)

	age	gender	impluse	pressurehight	pressurelow	glucose	class
1	64	male	66	160	83	160	negative
2	21	male	94	98	46	296	positive
3	55	male	64	-160	77	270	negative
4	64	male	70	120	55	270	positive
5	55	male	64	112	65	300	negative
6	58	female	61	112	58	87	negative
7	32	female	40	179	68	102	negative
8	63	male	60	214	82	87	positive
9	44	female	60	NA	81	135	negative
10	67		61	160	95	100	negative
11	NA	female	60	166	90	102	negative
12	63	female	60	150	10	198	negative
13	64	male	60	199	5	92	positive
14	54	female	94	122	67	97	negative
15	47	male	76	120	70	319	negative
16	61	male	81	NA	66	134	positive
17	86	female	73	114	68	87	positive
18	45	female	70	100	68	96	negative
19	37	female	72	107	86	274	negative
20	45	male	60	109	65	89	positive
21	60	male	92	151	78	301	negative
22	48	male	135	98	60	100	positive
23	52	male	76	109	85	227	positive
24	30	male	63	110	68	107	positive
25	NA	male	63	320	63	269	positive
26	72	male	64	106	68	111	positive
27	42	male	65	150	68	101	negative
28	72	female	64	325	60	95	negative
29	47	female	66	134	57	270	positive

Showing 1 to 29 of 150 entries, 7 total columns

Console

Data Quality

1. Data Correctness: Checks if there's any incorrect value.

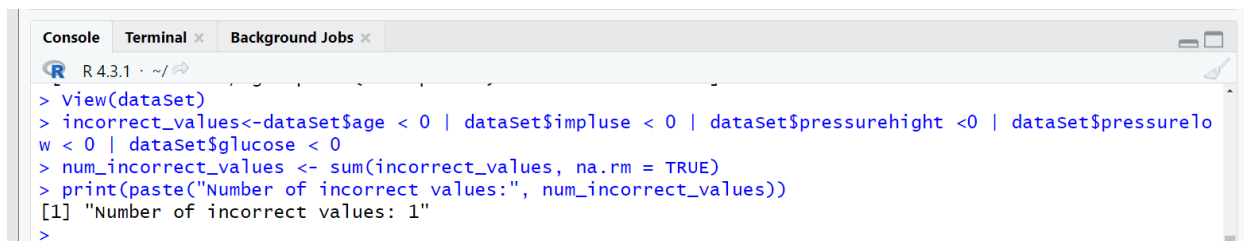
- Then summing up the occurrences of incorrect values.
- Printing the number of incorrect values.

Code:

```
incorrect_values<-dataSet$age < 0 | dataSet$impluse < 0 | dataSet$pressurehight <0 |  
dataSet$pressurelow < 0 | dataSet$glucose < 0
```

```
num_incorrect_values <- sum(incorrect_values, na.rm = TRUE)
```

```
print(paste("Number of incorrect values:", num_incorrect_values))
```



```
R 4.3.1 ~/  
> view(dataSet)  
> incorrect_values<-dataSet$age < 0 | dataSet$impluse < 0 | dataSet$pressurehight <0 | dataSet$pressurelow  
w < 0 | dataSet$glucose < 0  
> num_incorrect_values <- sum(incorrect_values, na.rm = TRUE)  
> print(paste("Number of incorrect values:", num_incorrect_values))  
[1] "Number of incorrect values: 1"  
>
```

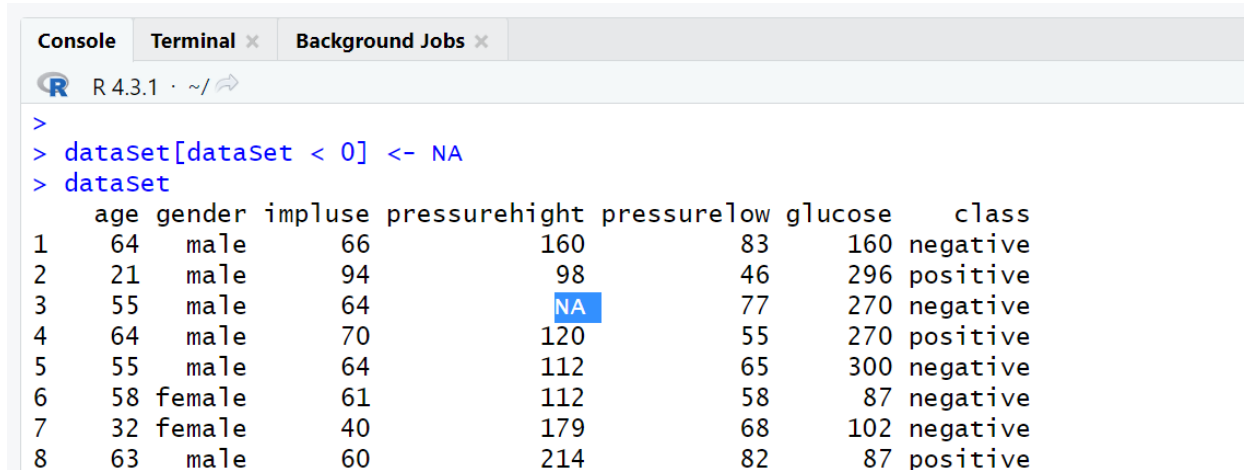
(Invalid Value)

Then, Replacing all negative/ incorrect values with 'NA'

Code:

```
dataSet[dataSet < 0] <- NA
```

dataSet



```
R 4.3.1 ~/  
>  
> dataSet[dataSet < 0] <- NA  
> dataSet  
  age gender impluse pressurehight pressurelow glucose  class  
1  64  male     66           160           83    160 negative  
2  21  male     94           98           46    296 positive  
3  55  male     64           NA           77    270 negative  
4  64  male     70          120           55    270 positive  
5  55  male     64          112           65    300 negative  
6  58 female     61          112           58     87 negative  
7  32 female     40          179           68    102 negative  
8  63  male     60          214           82     87 positive
```

2. Data Completeness: Checks missing values in dataset.

Code:

dataSet

is.na(dataSet)

```
Console Terminal x Background Jobs x
R 4.3.1 ~ /
> is.na(dataSet)
      age gender impluse pressurehigh pressurelow glucose class
[1,] FALSE  FALSE  FALSE          FALSE          FALSE  FALSE  FALSE
[2,] FALSE  FALSE  FALSE          FALSE          FALSE  FALSE  FALSE
[3,] FALSE  FALSE  FALSE           TRUE          FALSE  FALSE  FALSE
[4,] FALSE  FALSE  FALSE          FALSE          FALSE  FALSE  FALSE
[5,] FALSE  FALSE  FALSE          FALSE          FALSE  FALSE  FALSE
[6,] FALSE  FALSE  FALSE          FALSE          FALSE  FALSE  FALSE
[7,] FALSE  FALSE  FALSE          FALSE          FALSE  FALSE  FALSE
[8,] FALSE  FALSE  FALSE          FALSE          FALSE  FALSE  FALSE
[9,] FALSE  FALSE  FALSE           TRUE          FALSE  FALSE  FALSE
[10,] FALSE  TRUE  FALSE          FALSE          FALSE  FALSE  FALSE
[11,]  TRUE  FALSE  FALSE          FALSE          FALSE  FALSE  FALSE
[12,] FALSE  FALSE  FALSE          FALSE          FALSE  FALSE  FALSE
[13,] FALSE  FALSE  FALSE          FALSE          FALSE  FALSE  FALSE
```

Then,

Checking and Summing the total number of missing values.

```
> missing_values <- is.na(dataSet)
> num_missing_values <- sum(missing_values)
> print(paste("Number of missing values:", num_missing_values))
[1] "Number of missing values: 11"
> |
```

Missing Value

Discard Instances: Remove row with any missing values.

Code:

```
cleaned_data<- na.omit(dataSet)
```

cleaned_data

Console	Terminal x	Background Jobs x
R 4.3.1 · ~/		
> cleaned_data<- na.omit(dataSet)		
> cleaned_data		
	age	gender
1	64	male
2	21	male
4	64	male
5	55	male
6	58	female
7	32	female
8	63	male
12	63	female
13	64	male
14	54	female

Replace by Most Frequent/Average Value: This method generally refers to imputing missing values in a dataset with either 'the most frequent value (mode) for categorical data' or 'the average value (mean) for numerical data'.

So, we have findout mean values for every column. (age, gender, impulse, pressurehigh, pressurelow, glucose, class)

1. MEAN VALUE OF AGE: This code is used to calculate the mean value of the age.

Code:

```
mean_age_value <- mean(dataSet$age, na.rm = TRUE)
```

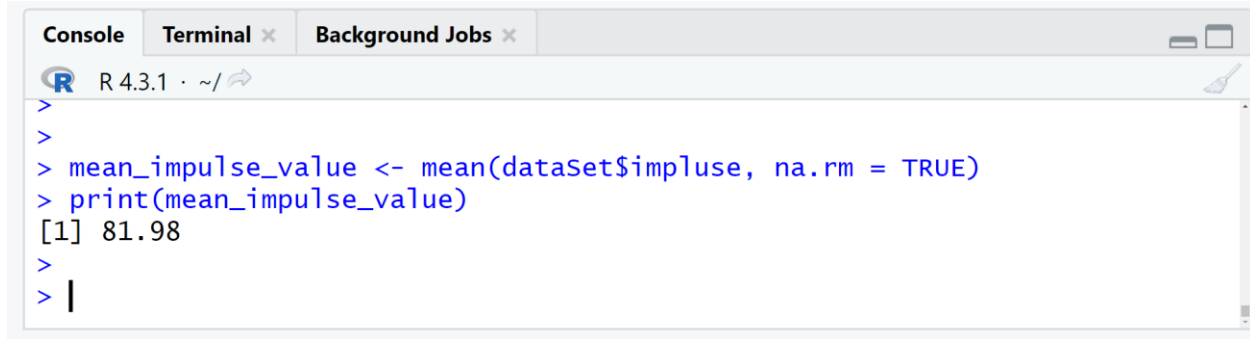
```
print(mean_age_value)
```

Console	Terminal x	Background Jobs x
R 4.3.1 · ~/		
>		
> mean_age_value <- mean(dataSet\$age, na.rm = TRUE)		
> print(mean_age_value)		
[1] 56.13793		
> print(mean_age_value)		
[1] 56.13793		
>		

2. MEAN VALUE OF IMPULSE: This code is used to calculate the mean value of the impulse.

Code:

```
mean_impulse_value <- mean(dataSet$impluse, na.rm = TRUE)
print(mean_impulse_value)
```

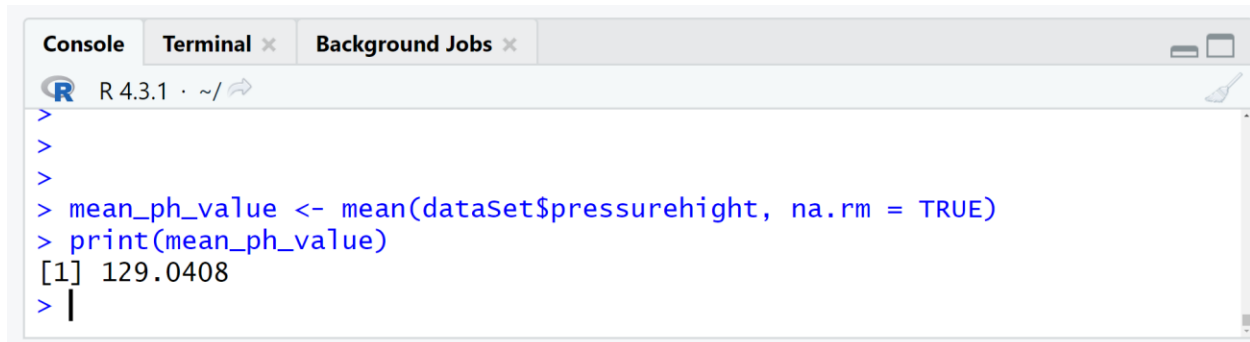
A screenshot of an R console window. The window has tabs for 'Console', 'Terminal', and 'Background Jobs'. The console shows the following code and output:

```
>
>
> mean_impulse_value <- mean(dataSet$impluse, na.rm = TRUE)
> print(mean_impulse_value)
[1] 81.98
>
> |
```

3. MEAN VALUE OF PRESSURE HIGHT: This code is used to calculate the mean value of the Pressureheight.

Code:

```
mean_ph_value <- mean(dataSet$pressurehight, na.rm = TRUE)
print(mean_ph_value)
```

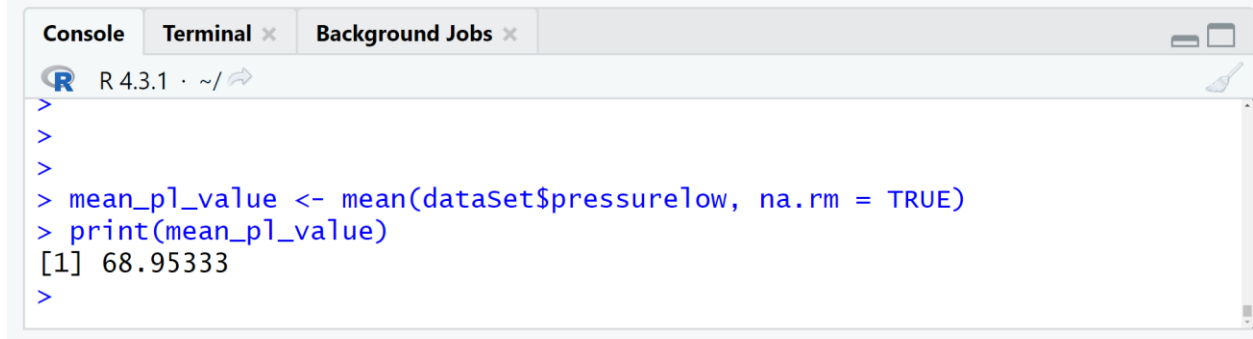
A screenshot of an R console window. The window has tabs for 'Console', 'Terminal', and 'Background Jobs'. The console shows the following code and output:

```
>
>
>
> mean_ph_value <- mean(dataSet$pressurehight, na.rm = TRUE)
> print(mean_ph_value)
[1] 129.0408
>
> |
```

4. **MEAN VALUE OF PRESSURE LOW:** This code is used to calculate the mean value of the pressurelow.

Code:

```
mean_pl_value <- mean(dataSet$pressurelow, na.rm = TRUE)
print(mean_pl_value)
```

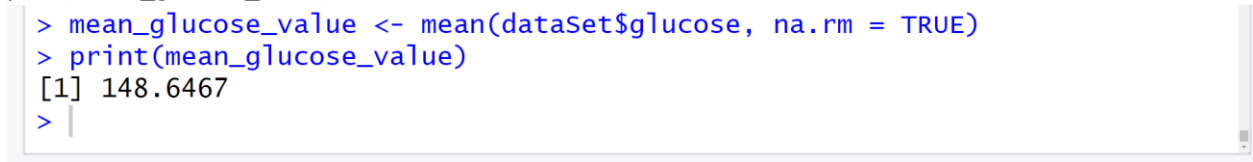
A screenshot of an R console window. The window has three tabs: 'Console', 'Terminal', and 'Background Jobs'. The 'Console' tab is active. The title bar shows 'R 4.3.1 · ~/'. The console output shows the following: a prompt '>' followed by three blank lines, then the command 'mean_pl_value <- mean(dataSet\$pressurelow, na.rm = TRUE)', followed by 'print(mean_pl_value)', and the output '[1] 68.95333'. The prompt '>' is visible at the end of the last line.

```
R 4.3.1 · ~/
>
>
>
> mean_pl_value <- mean(dataSet$pressurelow, na.rm = TRUE)
> print(mean_pl_value)
[1] 68.95333
>
```

5. **MEAN VALUE OF GLUCOSE:** This code is used to calculate the mean value of the glucose.

Code:

```
mean_glucose_value <- mean(dataSet$glucose, na.rm = TRUE)
print(mean_glucose_value)
> mean_glucose_value <- mean(dataSet$glucose, na.rm = TRUE)
> print(mean_glucose_value)
[1] 148.6467
>
```

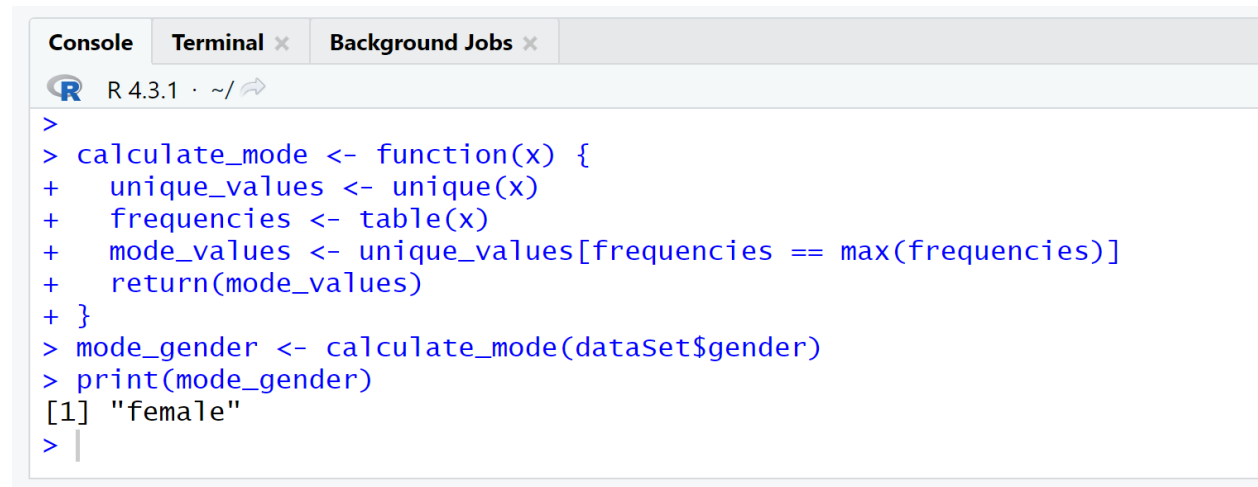
A screenshot of an R console window showing the execution of code to calculate the mean of glucose. The window has three tabs: 'Console', 'Terminal', and 'Background Jobs'. The 'Console' tab is active. The title bar shows 'R 4.3.1 · ~/'. The console output shows the following: a prompt '>' followed by two blank lines, then the command 'mean_glucose_value <- mean(dataSet\$glucose, na.rm = TRUE)', followed by 'print(mean_glucose_value)', and the output '[1] 148.6467'. The prompt '>' is visible at the end of the last line.

```
R 4.3.1 · ~/
>
>
> mean_glucose_value <- mean(dataSet$glucose, na.rm = TRUE)
> print(mean_glucose_value)
[1] 148.6467
>
```


6. MODE VALUE OF GENDER: (Categorical) This code is used to calculate the mode value of Gender column.

Code:

```
calculate_mode <- function(x) {  
  unique_values <- unique(x)  
  frequencies <- table(x)  
  mode_values <- unique_values[frequencies == max(frequencies)]  
  return(mode_values) }  
  
mode_gender <- calculate_mode(dataSet$gender)  
  
print(mode_gender)
```



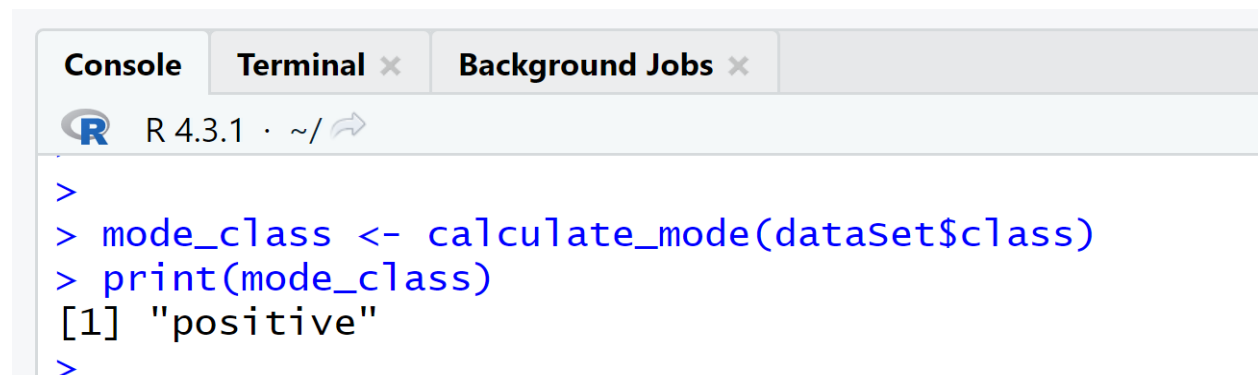
The screenshot shows an R console window with tabs for 'Console', 'Terminal', and 'Background Jobs'. The console output shows the definition of the 'calculate_mode' function and its application to the 'gender' column of a dataset, resulting in the output 'female'.

```
R 4.3.1 · ~/   
>  
> calculate_mode <- function(x) {  
+   unique_values <- unique(x)  
+   frequencies <- table(x)  
+   mode_values <- unique_values[frequencies == max(frequencies)]  
+   return(mode_values)  
+ }  
> mode_gender <- calculate_mode(dataSet$gender)  
> print(mode_gender)  
[1] "female"  
> |
```


7. MEAN VALUE OF CLASS: (Categorical) This code is used to calculate the mode value of the Class.

Code:

```
mode_class <- calculate_mode(dataSet$class)  
  
print(mode_class)
```



The screenshot shows an R console window with tabs for 'Console', 'Terminal', and 'Background Jobs'. The console output shows the application of the 'calculate_mode' function to the 'class' column of a dataset, resulting in the output 'positive'.

```
R 4.3.1 · ~/   
>  
> mode_class <- calculate_mode(dataSet$class)  
> print(mode_class)  
[1] "positive"  
>
```

Fixing Missing values

Replacing the missing values in every column with the average mean values (numerical) & mode values (categorical).

1. Fixing Age Column: There is no missing values in age column, but with the process and code we can fix all the column.

Code:

```
dataSet$age <- round(dataSet$age, 0)
```

```
dataSet[is.na(dataSet$age),"age"] <- mean_age_value
```

```
print(dataSet)
```

Console Terminal x Background Jobs x

R 4.3.1 · ~/

```
> dataSet$age <- round(dataSet$age, 0)
> dataSet[is.na(dataSet$age),"age"] <- mean_age_value
> print(dataSet)
```

	age	gender	impluse	pressurehigh	pressurelow	glucose	class
1	64	male	66	160	83	160	negative
2	21	male	94	98	46	296	positive
3	55	male	64	NA	77	270	negative
4	64	male	70	120	55	270	positive
5	55	male	64	112	65	300	negative
6	58	female	61	112	58	87	negative
7	32	female	40	179	68	102	negative
8	63	male	60	214	82	87	positive
9	44	female	60	NA	81	135	negative
10	67	<NA>	61	160	95	100	negative
11	56	female	60	166	90	102	negative
12	63	female	60	150	10	198	negative
13	64	male	60	199	5	92	positive
14	54	female	94	122	67	97	negative
15	47	male	76	120	70	319	negative
16	61	male	81	NA	66	134	positive
17	86	female	73	114	68	87	positive
18	45	female	70	100	68	96	negative
19	37	female	72	107	86	274	negative
20	45	male	60	109	65	89	positive
21	60	male	92	151	78	301	negative

2. Fixing Impulse Column:

Code:

```
dataSet$impluse <- round(dataSet$impluse, 2)
dataSet[is.na(dataSet$impluse),"impluse"] <- mean_impluse_value
print(dataSet)
```

3. Fixing Pressure Height Column:

Code:

```
dataSet$pressureheight <- round(dataSet$pressureheight, 2)
dataSet[is.na(dataSet$pressureheight),"pressureheight"] <- mean_ph_value
print(dataSet)
```

4. Fixing Pressure Low Column:

Code:

```
dataSet$pressurelow <- round(dataSet$pressurelow, 1)
dataSet[is.na(dataSet$pressurelow),"pressurelow"] <- mean_pl_value
print(dataSet)
```

5. Fixing GLUCOSE Column:

Code:

```
dataSet$glucose <- round(dataSet$glucose, 1)
dataSet[is.na(dataSet$glucose),"glucose"] <- mean_glucose_value
print(dataSet)
```

6. Fixing Gender Column(mode):

Code:

```
dataSet[is.na(dataSet$gender),"gender"] <- mode_gender
print(dataSet)
```

7. Fixing Class Column(mode):

Code:

```
dataSet[is.na(dataSet$class),"class"] <- mode_class
```

```
print(dataSet)
```

Here, The Fixed Dataset

	age	gender	impluse	pressure	height	pressurelow	glucose	class
1	64.00000	male	66	160.0000		83	160	negative
2	21.00000	male	94	98.0000		46	296	positive
3	55.00000	male	64	129.0408		77	270	negative
4	64.00000	male	70	120.0000		55	270	positive
5	55.00000	male	64	112.0000		65	300	negative
6	58.00000	female	61	112.0000		58	87	negative
7	32.00000	female	40	179.0000		68	102	negative
8	63.00000	male	60	214.0000		82	87	positive
9	44.00000	female	60	129.0408		81	135	negative
10	67.00000	female	61	160.0000		95	100	negative
11	56.13793	female	60	166.0000		90	102	negative
12	63.00000	female	60	150.0000		10	198	negative
13	64.00000	male	60	199.0000		5	92	positive
14	54.00000	female	94	122.0000		67	97	negative
15	47.00000	male	76	120.0000		70	319	negative
16	61.00000	male	81	129.0408		66	134	positive
17	86.00000	female	73	114.0000		68	87	positive
18	45.00000	female	70	100.0000		68	96	negative
19	37.00000	female	72	107.0000		86	274	negative
20	45.00000	male	60	109.0000		65	89	positive
21	60.00000	male	92	151.0000		78	301	negative
22	48.00000	male	135	98.0000		60	100	positive
23	52.00000	male	76	109.0000		85	227	positive
24	30.00000	male	63	110.0000		68	107	positive
25	56.13793	male	63	320.0000		63	269	positive
26	72.00000	male	64	106.0000		68	111	positive
27	42.00000	male	65	150.0000		68	101	negative
28	72.00000	female	64	325.0000		60	95	negative
29	47.00000	female	66	134.0000		57	279	positive
30	63.00000	male	66	135.0000		55	166	negative
31	54.00000	male	125	131.0000		82	95	positive
32	35.00000	male	62	137.0000		61	321	negative
33	68.00000	male	61	121.0000		49	98	positive
34	56.00000	female	60	145.0000		62	105	negative
35	50.00000	male	61	136.0000		70	136	positive
36	64.00000	male	58	156.0000		76	82	positive
37	56.13793	male	60	166.0000		82	117	negative
38	64.00000	male	65	155.0000		75	107	negative

The dataset shows there's no missing values

Outliers

Setting the layout to 2 rows and 2 columns, so that its easy to show the plots together.

Code:

```
par(mfrow = c(2, 2))
```

Now Box plot for (Age, Impluse, Pressureheight, Pressurelow, Glucose)

Code:

```
boxplot(dataSet$age, main = "Age")
```

```
boxplot(dataSet$impluse, main = "Impluse")
```

```
boxplot(dataSet$pressurehight, main = "Pressure Hight")
```

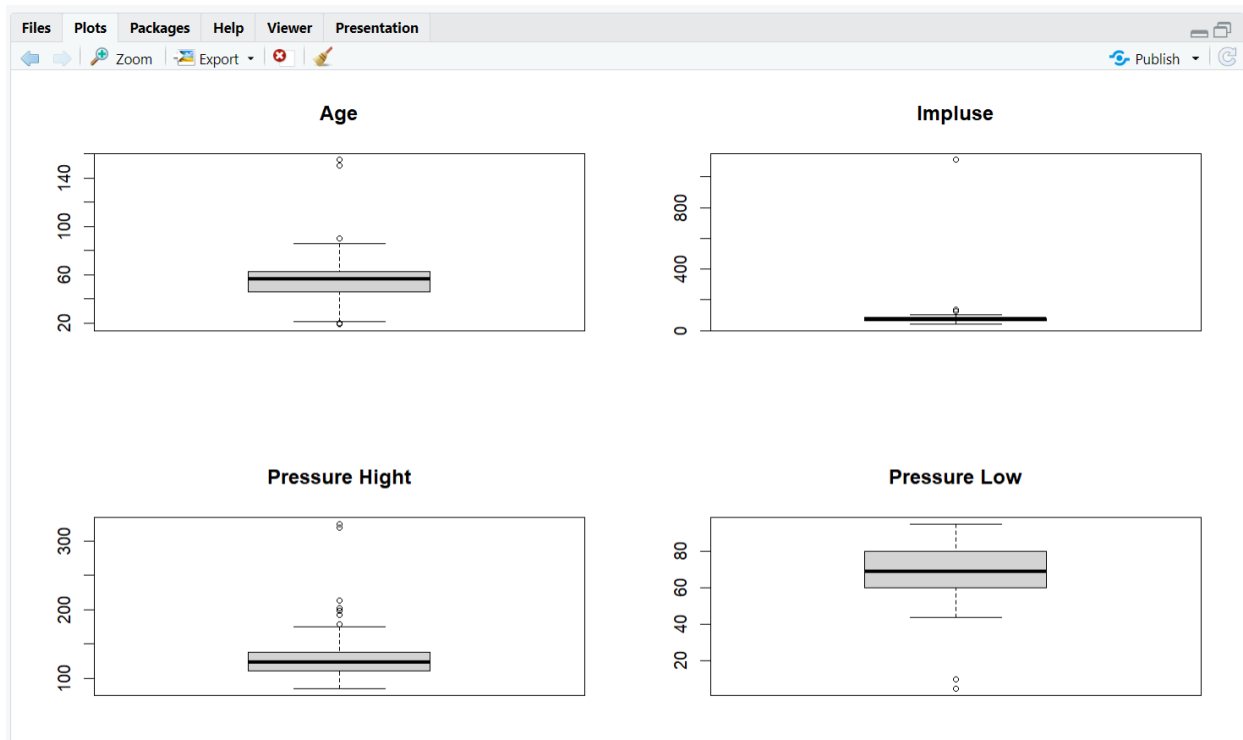
```
boxplot(dataSet$pressurelow, main = "Pressure Low")
```

```
boxplot(dataSet$glucose, main = "Glucose")
```

Bar plot for Categorical:

```
barplot(table(dataSet$gender), main = "Gender Distribution", xlab = "Gender", ylab = "Count")
```

```
barplot(table(dataSet$class), main = "Class Distribution", xlab = "Class", ylab = "Count")
```





Box Plot Info:

Age: It is shown that most of the values are in the range of 20 to 80 and the noisy values are above 140.

Impulse: It can be seen that most of the values are in the range of 40 to 135 and the outlier is 1111.

Pressure high: It is shown that most of the values are in the range of 85 to 215 and the noisy values are above 300.

Pressure low: It is shown that most of the values are in the range of 40 to 95 and the noisy values are below 15.

Glucose: It is shown that most of the values are in the range of 60 to 300 and the noisy values are above 300.

Bar Plot Info:

Gender: It is shown that the count of female is 50 and male count is around 100.

Class: It is shown that the count of negative is 60 and positive count is around 100.

Data Type and Conversion

Converting Categorical data in numerical data. [gender(1,2) & class(3,4)]

Code:

```
dataSet$gender<-factor(dataSet$gender,  
  levels = c("male","female"),  
  labels = c(1,2))
```

dataSet

```
dataSet$class<-factor(dataSet$class,  
  levels = c("positive","negative"),  
  labels = c(3,4))
```

dataSet

Console Terminal × Background Jobs ×

R 4.3.1 · ~/ ↻

	age	gender	im	pluse	pressure	hight	pressure	low	glucose	class
1	64.00000	1	66	160.0000	83	160	4			
2	21.00000	1	94	98.0000	46	296	3			
3	55.00000	1	64	129.0408	77	270	4			
4	64.00000	1	70	120.0000	55	270	3			
5	55.00000	1	64	112.0000	65	300	4			
6	58.00000	2	61	112.0000	58	87	4			
7	32.00000	2	40	179.0000	68	102	4			
8	63.00000	1	60	214.0000	82	87	3			
9	44.00000	2	60	129.0408	81	135	4			
10	67.00000	2	61	160.0000	95	100	4			
11	56.13793	2	60	166.0000	90	102	4			
12	63.00000	2	60	150.0000	10	198	4			
13	64.00000	1	60	199.0000	5	92	3			
14	54.00000	2	94	122.0000	67	97	4			
15	47.00000	1	76	120.0000	70	319	4			
16	61.00000	1	81	129.0408	66	134	3			
17	86.00000	2	73	114.0000	68	87	3			
18	45.00000	2	70	100.0000	68	96	4			
19	37.00000	2	72	107.0000	86	274	4			
20	45.00000	1	60	109.0000	65	89	3			
21	60.00000	1	92	151.0000	78	301	4			
22	48.00000	1	135	98.0000	60	100	3			
23	52.00000	1	76	109.0000	85	227	3			
24	30.00000	1	63	110.0000	68	107	3			

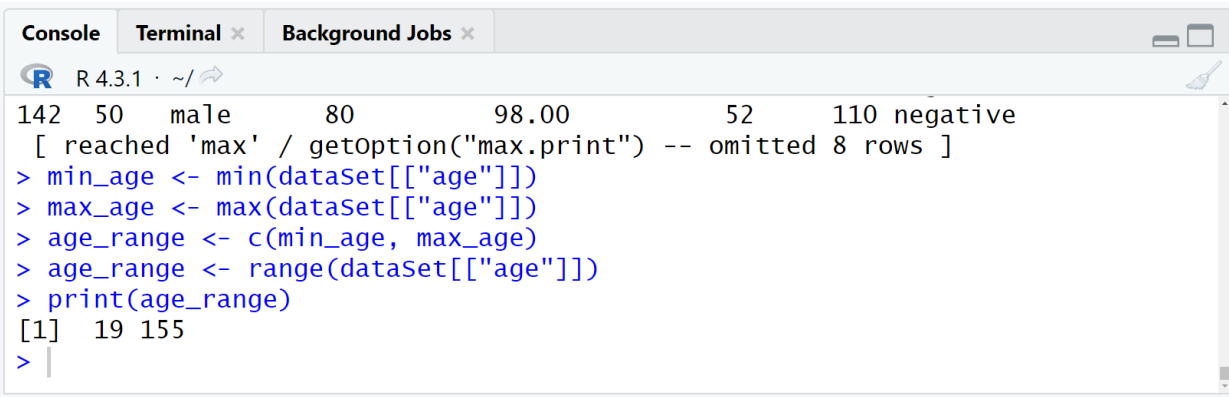
Data Exploration

Measure of Spread: (Range)


1.AGE RANGE

Code:

```
min_age <- min(dataSet[["age"]])  
  
max_age <- max(dataSet[["age"]])  
  
age_range <- c(min_age, max_age)  
  
age_range <- range(dataSet[["age"]])  
  
print(age_range)
```



The screenshot shows an R console window with the following content:

```
R 4.3.1 · ~/   
142 50 male 80 98.00 52 110 negative  
[ reached 'max' / getOption("max.print") -- omitted 8 rows ]  
> min_age <- min(dataSet[["age"]])  
> max_age <- max(dataSet[["age"]])  
> age_range <- c(min_age, max_age)  
> age_range <- range(dataSet[["age"]])  
> print(age_range)  
[1] 19 155  
> |
```

This code is used to find the range of age.

2.IMPLUSE RANGE

Code:

```
min_impluse <- min(dataSet[["impluse"]])  
  
max_impluse <- max(dataSet[["impluse"]])  
  
impulse_range <- c(min_impluse, max_impluse)  
  
impulse_range <- range(dataSet[["impluse"]])  
  
print(impulse_range)
```

```
Console Terminal x Background Jobs x
R 4.3.1 · ~/
>
> min_impluse <- min(dataSet[["impluse"]])
> max_impluse <- max(dataSet[["impluse"]])
> impulse_range <- c(min_impluse, max_impluse)
> impulse_range <- c(min_impluse, max_impluse)
> impulse_range <- c(min_impluse, max_impluse)
> impulse_range <- range(dataSet[["impluse"]])
> print(impulse_range)
[1] 40 1111
> |
```

This code is used to find the range of impulse

Variance & Standard Deviation: (Age)

Variance Code:

```
age_variance <- var(dataSet[["age"]])
```

```
print(age_variance)
```

Standard Deviation Code:

```
age_std_dev <- sd(dataSet[["age"]])
```

```
print(age_std_dev)
```

```
Console Terminal x Background Jobs x
R 4.3.1 · ~/
>
> age_variance <- var(dataSet[["age"]])
> print(age_variance)
[1] 288.9083
> age_std_dev <- sd(dataSet[["age"]])
> print(age_std_dev)
[1] 16.9973
> |
```

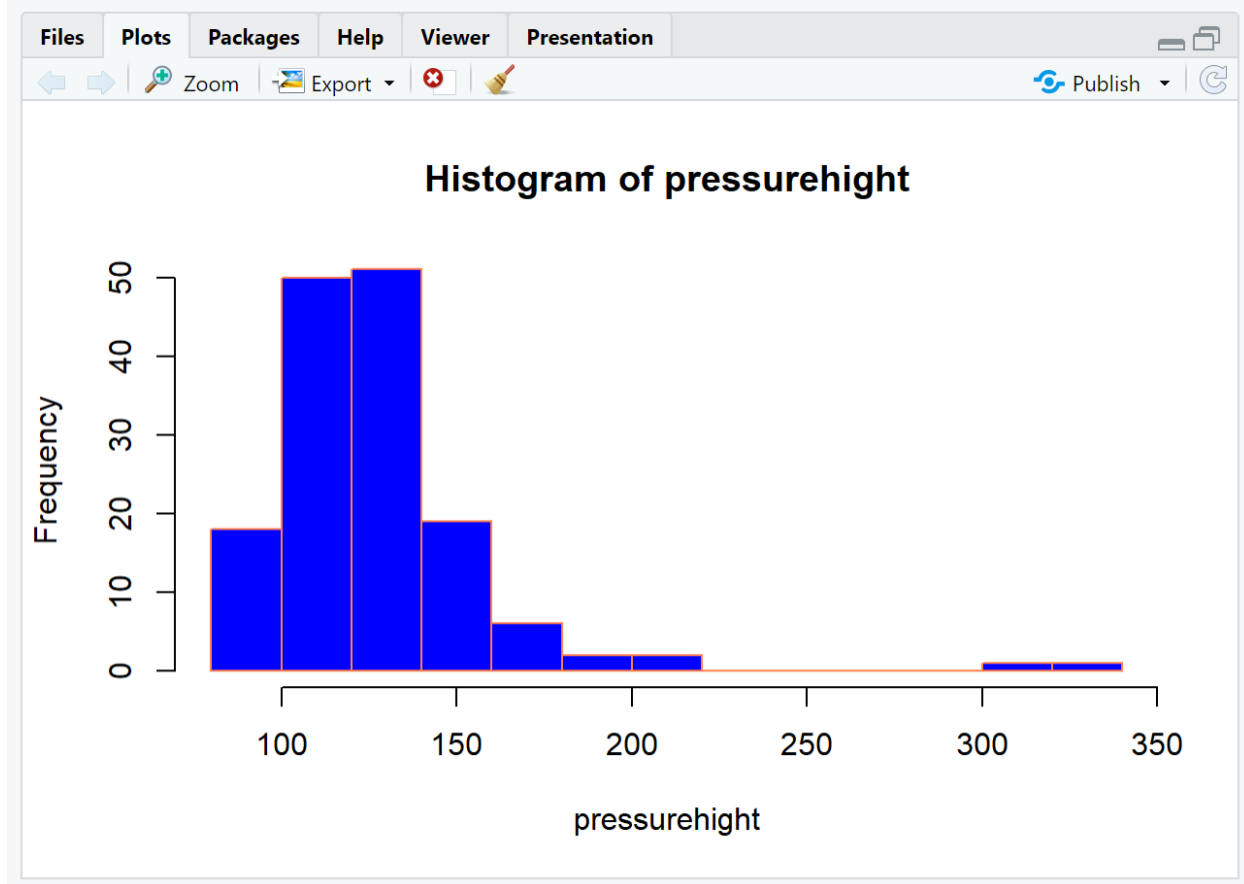
This code is used to find the variance & standard deviation of Age

Univariate Visualization

Histogram: (Pressurehight)

Code:

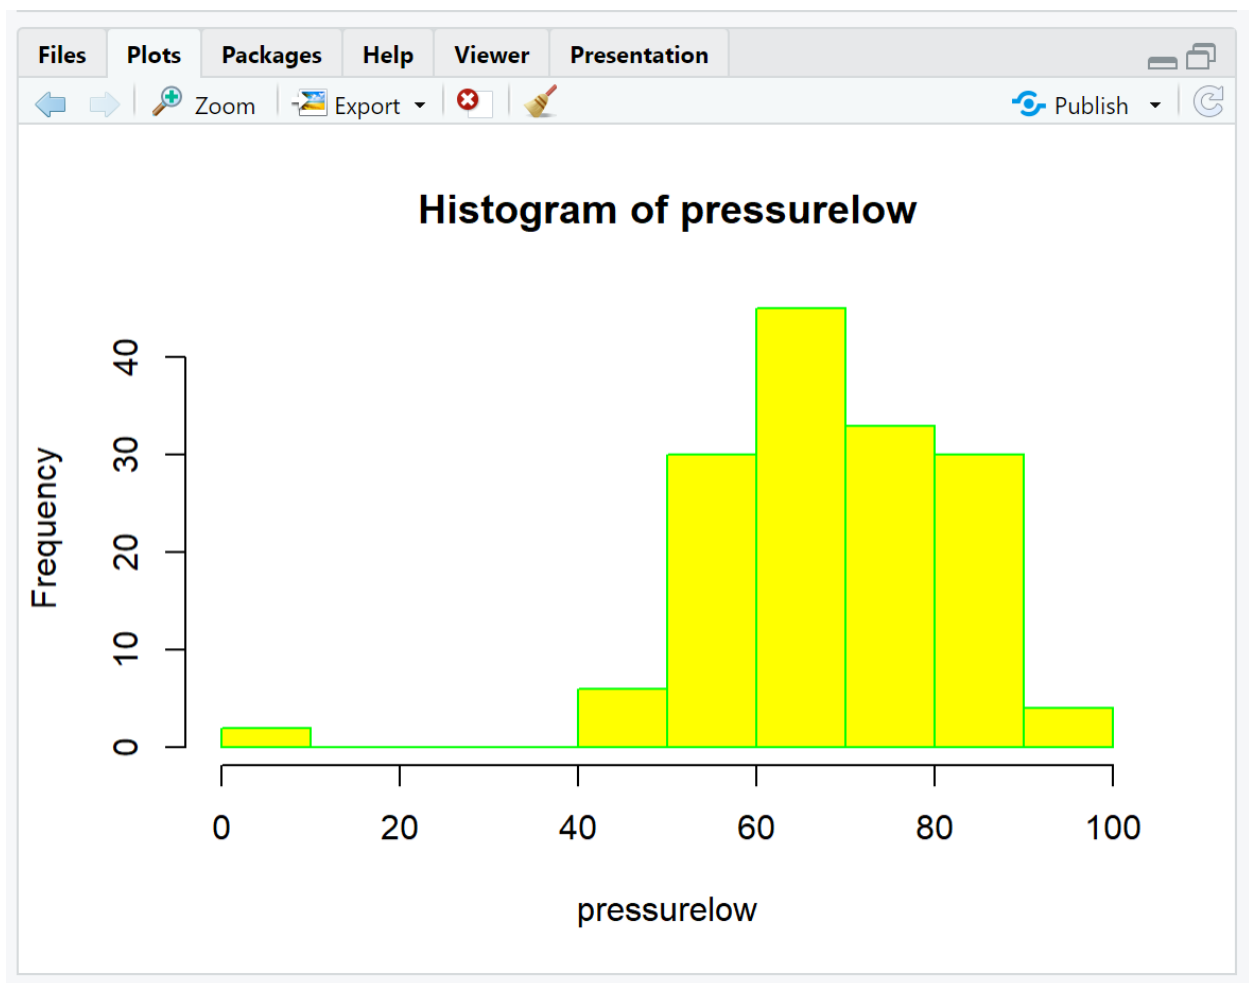
```
hist(dataSet[["pressurehight"]],  
      main= paste("Histogram of", "pressurehight"),  
      xlab = "pressurehight",  
      ylab = "Frequency",  
      col = "blue",  
      border = "coral")
```



For, Pressurelow

Code:

```
hist(dataSet[["pressurelow"]],  
      main= paste("Histogram of", "pressurelow"),  
      xlab = "pressurelow",  
      ylab = "Frequency",  
      col = "yellow",  
      border = "green")
```



Bar Plot: (already done in Outliers)